

SABANCI UNIVERSITY
CS406/531: Parallel Computing
Spring 2018-2019, Midterm I
Duration: 110 minutes
Grade: / 100

Q1 (15pts): Suppose a fraction r of the runtime of a serial program is “*perfectly parallelized*,” and the remaining fraction $1 - r$ is “*inherently serial*.” Give an upper bound on the speedup one can have.

Q2 (15 pts): Describe the following terms and their differences in at most four sentences each.

- a) (5 pts) Latency **and** bandwidth
- b) (5 pts) Spatial locality **and** temporal locality
- c) (5 pts) Shared address space **and** distributed address space

Q3 (10 pts): Implement a parallel function with OpenMP to reverse an array in five lines. Try to be as efficient as possible. The function **must be not more than 5 lines and be an in-place function**, i.e., using extra memory is not allowed.

```
#include <iostream>
#include <memory>

using namespace std;

//5 lines including pragma – no memory allocation
void reverse(int* arr, int length) {


}

int main() {
    int N = 1000000;
    int* arr = new int[N];
    for(int i = 0; i < N; i++) {
        arr[i] = rand();
    }
    reverse(arr, N);
}
```


Q4 (10pts): Answer the following questions about the charts below: (**E** is efficiency, **p** is number of processors, **W** is the problem size).

- Which chart (left or right) is correct for **all** parallel systems (architecture + algorithm)? Why?
- Which chart (left or right) is correct for **some** parallel systems (architecture + algorithm)? What do we call such systems and why?


Answers without correct explanations won't get partial credit.



Q5 (15pts): Given two sample task dependency graphs as in the figures (5.1) and (5.2) of **N uniform** tasks. **Assume the graphs below are only instances of a family of task dependency graphs following the same pattern**, so **n** and **N** can be different.



(5.1) $N = n^2$



(5.2) $N = 2^n - 1$

- (5pts) What is the critical path length for each graph **in terms of n and N**?
- (5pts) What is the degree of concurrency for each graph **in terms of n and N**?
- (5pts) What is the speedup and efficiency one can obtain **in terms of n and N** if the number of processing elements is equal to the degree of concurrency? Show all your work.

Q6 (15pts): Consider a processor operating in 4GHz connected to a DRAM with a latency of 25ns (no caches). Assume the processor has two multiply add units and is capable of executing four instructions at each cycle of 1ns.

- a. (3pts) What is the peak performance of the processor in FLOPS?
- b. (5pts) Consider the problem of computing the dot product of two vectors (for one multiply-add, we need two data items): What is the peak performance in FLOPS for that algorithm? Explain the math you use.
- c. (3pts) Assume that you have a single cycle cache and the block size (memory bandwidth) is one word. What is the peak performance for dot product computation in this case?
- d. (4pts) What is the peak performance again with a single cycle cache if the block size is 8?

Q7 (10 pts): Suppose that you implemented a function **party()** as shown below:

```
int party() {
    int a = 0;
#pragma omp parallel num_threads(3)
{
    for(int i = 0; i < 5; i++) {
        a = a + 1;
    }
}
return a;
```

- a) (3 pts) While building the executable, you forgot to use **-fopenmp**. The code still compiles; what are the maximum and minimum values the function party can return?
- b) (7 pts) You added **-fopenmp** to build the executable. What are the maximum and minimum possible values the function can return?

Q8 (10 pts): Please answer the following:

- a. (5 pts) Given a sequential algorithm requiring **O(n)** operations on **n** numbers, assume that you devised a parallel version requiring **O(n/p + log p)** time with **p** processors. What is the efficiency (E) of your algorithm?
- b. (5 pts) (Continue from part (a)): When (up to how many processors) is your algorithm cost optimal? Show all your math, only the answer will not get partial credits.