

# Grids, Thread Blocks and Threads

## Grid

Thread Block 0, 0

0,0	0,1	0,2	0,3
1,0	1,1	1,2	2,3
2,0	2,1	2,2	2,3

Thread Block 0, 1

0,0	0,1	0,2	0,3
1,0	1,1	1,2	2,3
2,0	2,1	2,2	2,3

Thread Block 0, 2

0,0	0,1	0,2	0,3
1,0	1,1	1,2	2,3
2,0	2,1	2,2	2,3

Thread Block 1, 0

0,0	0,1	0,2	0,3
1,0	1,1	1,2	2,3
2,0	2,1	2,2	2,3

Thread Block 1, 1

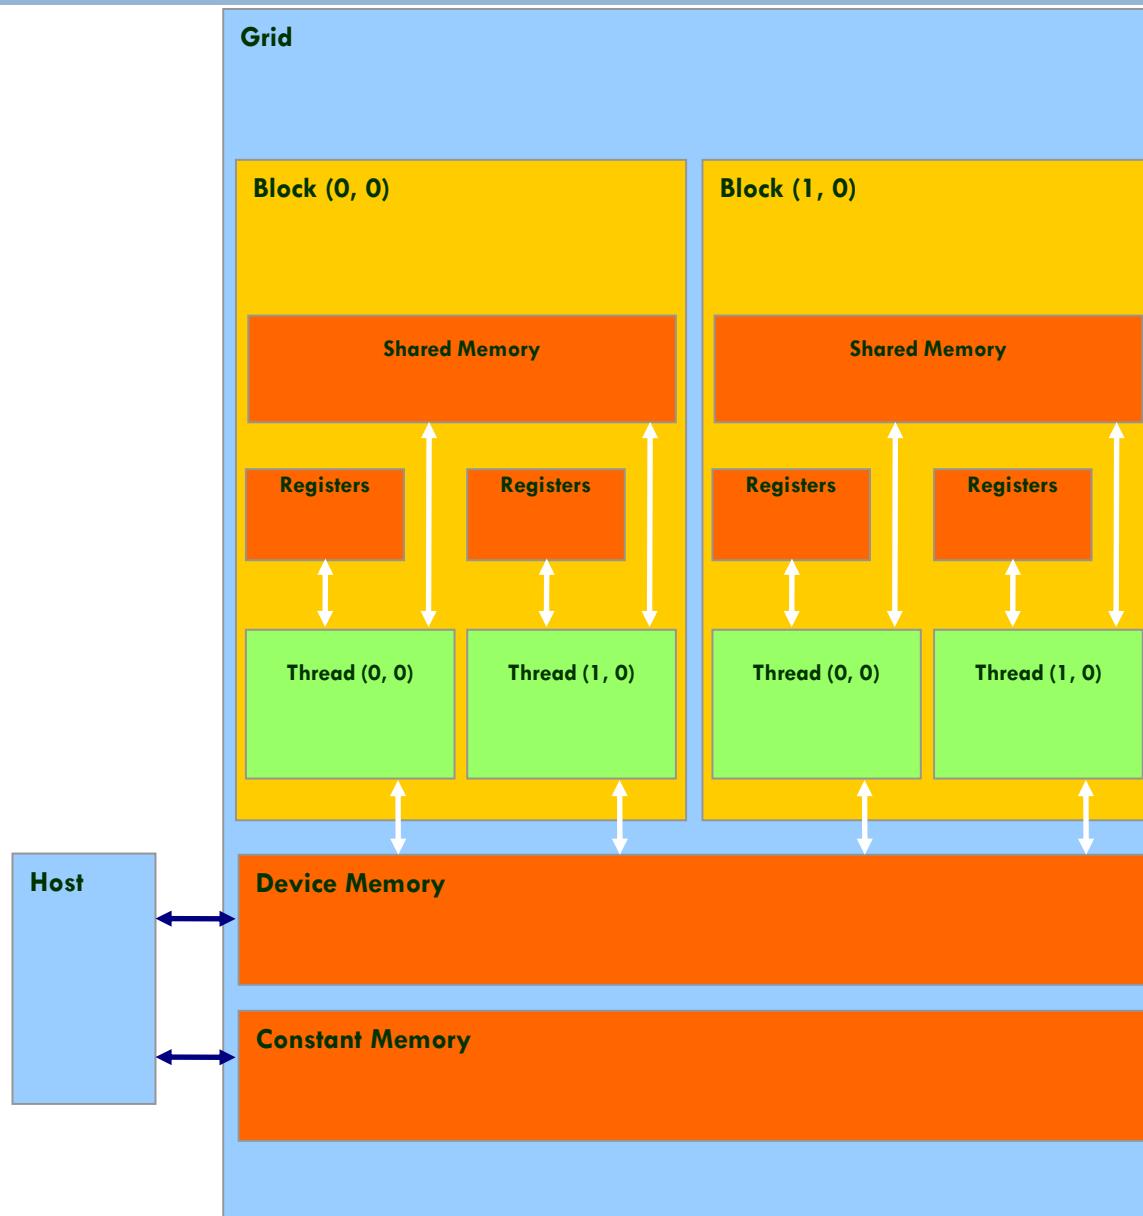
0,0	0,1	0,2	0,3
1,0	1,1	1,2	2,3
2,0	2,1	2,2	2,3

Thread Block 1, 2

0,0	0,1	0,2	0,3
1,0	1,1	1,2	2,3
2,0	2,1	2,2	2,3

# Hardware Memory Spaces in CUDA

4



# Vector addition GPU code

5

```
// compute vector sum c = a + b
// each thread performs one pair-wise addition
__global__ void vector_add(float* A, float* B, float* C) {
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    C[i] = A[i] + B[i];
}
```

GPU code

```
int main() {
    // initialization code here ...

    // launch N/256 blocks of 256 threads each
    vector_add<<< N/256, 256 >>>(deviceA, deviceB, deviceC);

    // cleanup code here ...
}
```

Host code

(can be in the same file)

- The host always initiates work for the device
  - For Kepler GPUs, device can generate work for itself (Dynamic Parallelism)
- Examples for common extensions are
  - `__global__` defines a GPU kernel that is callable by the CPU
  - `__device__` defines a GPU function that is callable by a GPU kernel or by another device function, but not callable by a CPU (host)
  - `__host__` or no qualifier marks a CPU function

- A CUDA kernel is a **grid** of **thread blocks**
  - The grid can be 1D, 2D or 3D
- A **thread block** is, in turn, a 1D, 2D, or 3D **array of threads**
- A thread is executed by exactly one **stream processor** or **CUDA core**
- A thread block is executed by exactly one **Streaming Multiprocessor (SM)**
- However, CUDA cores and SM can interleave execution of multiple threads and thread blocks

- CUDA maintains special variables that store:
  - thread index within a thread block `threadIdx(.x, .y, .z)`
  - block index within a kernel grid `blockIdx(.x, .y, .z)`
  - dimension of a thread block `blockDim(.x, .y, .z)`
  - dimension of a kernel grid `gridDim(.x, .y, .z)`
- These variables are used without declaration
- Mainly used in assigning the work per block/thread
- The maximum values of the (x, y, z) in both `blockDim` and `gridDim` are GPU-dependent
- At least the x component needs to be declared. The default value for the y and z components is 1

```
dim3 dimBlock(3)  
dim3 dimGrid(4,2)
```

(0,0)	(1,0)	(2,0)	(3,0)
0 1 2	0 1 2	0 1 2	0 1 2
0 1 2	0 1 2	0 1 2	0 1 2
(0,1)	(1,1)	(2,1)	(3,1)

```
dim3 dimBlock(4, 2)  
dim3 dimGrid(3)
```

(0)	(1)	(2)
0,0 1,0 2,0 3,0 0,1 1,1 2,1 3,1	0,0 1,0 2,0 3,0 0,1 1,1 2,1 3,1	0,0 1,0 2,0 3,0 0,1 1,1 2,1 3,1

```
dim3 dimBlock(4, 2)  
dim3 dimGrid(3, 3)
```

(0,0)	(1,0)	(2,0)
0,0 1,0 2,0 3,0 0,1 1,1 2,1 3,1	0,0 1,0 2,0 3,0 0,1 1,1 2,1 3,1	0,0 1,0 2,0 3,0 0,1 1,1 2,1 3,1
(0,1)	(1,1)	(2,1)
0,0 1,0 2,0 3,0 0,1 1,1 2,1 3,1	0,0 1,0 2,0 3,0 0,1 1,1 2,1 3,1	0,0 1,0 2,0 3,0 0,1 1,1 2,1 3,1
(0,2)	(1,2)	(2,2)
0,0 1,0 2,0 3,0 0,1 1,1 2,1 3,1	0,0 1,0 2,0 3,0 0,1 1,1 2,1 3,1	0,0 1,0 2,0 3,0 0,1 1,1 2,1 3,1

# Thread Scheduling

7

- ❑ Order in which thread blocks are scheduled is undefined!
  - ❑ any possible interleaving of blocks should be valid
  - ❑ presumed to run to completion without preemption
  - ❑ can run in any order
  - ❑ can run concurrently OR sequentially
- ❑ Order of threads within a block is also undefined!

# Global synchronization

8

- Q: How do we do global synchronization with these scheduling semantics?

# Global synchronization

9

- Q: How do we do global synchronization with these scheduling semantics?
- A1: Not possible!

# Global synchronization

10

- Q: How do we do global synchronization with these scheduling semantics?
- A1: Not possible!
- A2: Finish a grid, and start a new one!

# Global synchronization

11

- Q: How do we do global synchronization with these scheduling semantics?
- A1: Not possible!
- A2: Finish a grid, and start a new one!

```
step1<<<grid1,blk1>>>(...);  
// CUDA ensures that all writes from step1 are complete.  
step2<<<grid2,blk2>>>(...);
```

- We don't have to copy the data back and forth!

# Atomics

12

- Guarantee that only a single thread has access to a piece of memory during an operation
- No dropped data, but ordering is still arbitrary
- Different types of atomic instructions
- Atomic Add, Sub, Exch, Min, Max, Inc, Dec, CAS, And, Or, Xor
- Can be done on device memory and shared memory
- Much more expensive than load + operation + store

# Example: Histogram

13

```
// Determine frequency of colors in a picture.  
// Colors have already been converted into integers  
// between 0 and 255.  
// Each thread looks at one pixel,  
// and increments a counter  
  
__global__ void histogram(int* colors, int* buckets)  
{  
    int i = threadIdx.x + blockDim.x * blockIdx.x;  
    int c = colors[i];  
    buckets[c] += 1;  
}
```

# Example: Histogram

14

```
// Determine frequency of colors in a picture.  
// Colors have already been converted into integers  
// between 0 and 255.  
// Each thread looks at one pixel,  
// and increments a counter  
  
__global__ void histogram(int* colors, int* buckets)  
{  
    int i = threadIdx.x + blockIdx.x * blockIdx.x;  
    int c = colors[i];  
    buckets[c] += 1;  
}
```

*Why is this incorrect?*

# Example: Histogram

15

```
// Determine frequency of colors in a picture.  
// Colors have already been converted into integers  
// between 0 and 255.  
// Each thread looks at one pixel,  
// and increments a counter atomically  
  
__global__ void histogram(int* colors, int* buckets)  
{  
    int i = threadIdx.x + blockDim.x * blockIdx.x;  
    int c = colors[i];  
    atomicAdd(&buckets[c], 1);  
}
```

# Example: Work queue

16

```
// For algorithms where the amount of work per item
// is highly non-uniform, it often makes sense to
// continuously grab work from a queue.
```

global

```
void workq(int* work_q, int* q_counter,
           int queue_max, int* output)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    int q_index = atomicInc(q_counter, queue_max);
    int result = do_work(work_q[q_index]);
    output[i] = result;
}
```

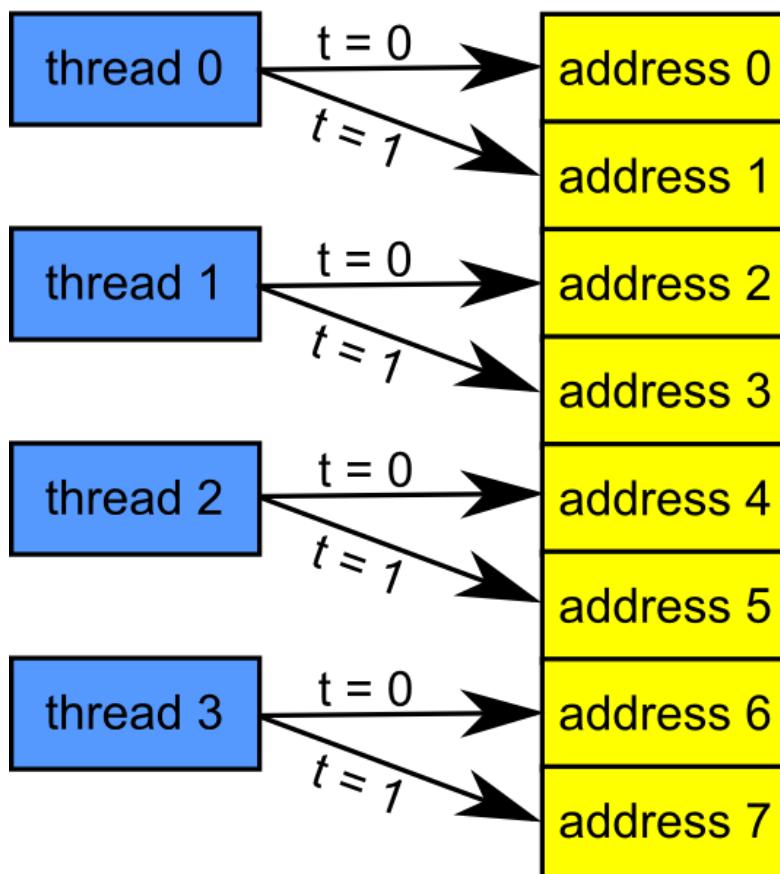
# CUDA: optimizing your application

## Coalescing

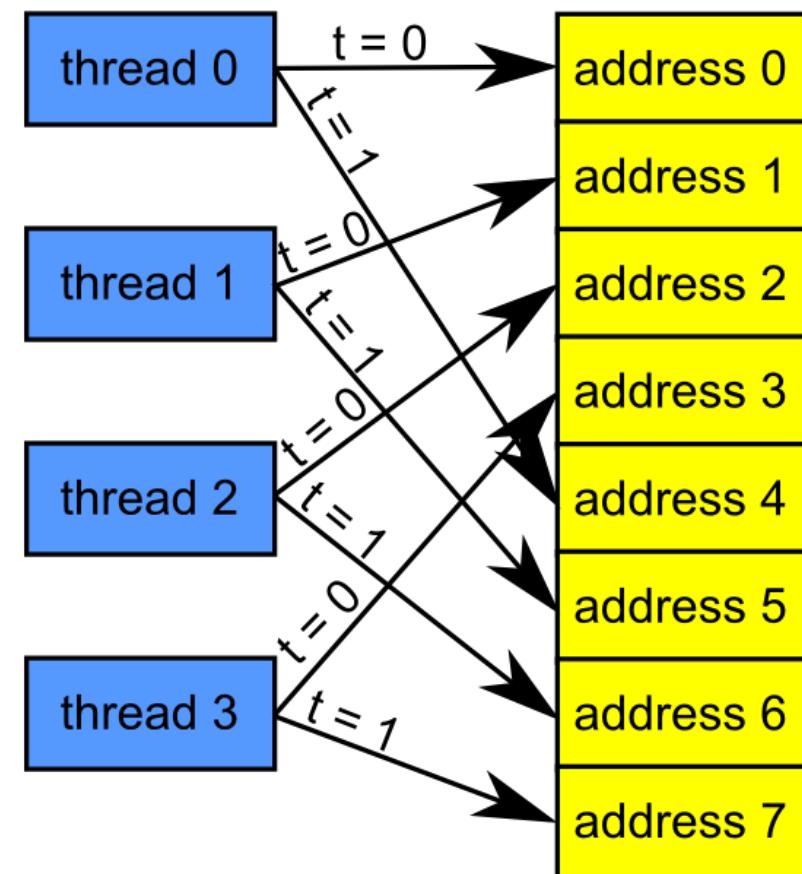
# Coalescing

18

traditional multi-core  
optimal memory access pattern



many-core GPU  
optimal memory access pattern



# Consider the stride of your accesses

19

```
__global__ void foo(int* input, float3* input2) {
    int i = blockDim.x * blockIdx.x + threadIdx.x;

    // Stride 1, OK!
    int a = input[i];

    // Stride 2, half the bandwidth is wasted
    int b = input[2*i];

    // Stride 3, 2/3 of the bandwidth wasted
    float c = input2[i].x;
}
```

# Example: Array of Structures (AoS)

20

```
struct record {  
    int key;  
    int value;  
    int flag;  
};  
  
record *d_records;  
cudaMalloc( (void**) &d_records, ... );
```

# Example: Structure of Arrays (SoA)

21

```
Struct SoA {  
    int* keys;  
    int* values;  
    int* flags;  
};  
  
SoA d_SoA_data;  
cudaMalloc( (void**) &d_SoA_data.keys, ... );  
cudaMalloc( (void**) &d_SoA_data.values, ... );  
cudaMalloc( (void**) &d_SoA_data.flags, ... );
```

# Example: SoA vs AoS

22

```
__global__ void bar(record* AoS_data,
                     SoA  SoA_data) {
    int i = blockDim.x * blockIdx.x + threadIdx.x;

    // AoS wastes bandwidth
    int key1 = AoS_data[i].key;

    // SoA efficient use of bandwidth
    int key2 = SoA_data.keys[i];
}
```

# Memory Coalescing

23

- Structure of arrays is often better than array of structures
- Very clear win on regular, stride 1 access patterns
- Unpredictable or irregular access patterns are case-by-case
- Can lose a factor of 10 – 30!

# CUDA: optimizing your application

## Shared Memory

# Using shared memory

25

```
// Adjacent Difference application:  
// compute result[i] = input[i] - input[i-1]  
  
__global__ void adj_diff_naive(int *result, int *input) {  
    // compute this thread's global index  
    unsigned int i = blockDim.x * blockIdx.x + threadIdx.x;  
  
    if(i > 0) {  
        // each thread loads two elements from device memory  
        int x_i = input[i];  
        int x_i_minus_one = input[i-1];  
  
        result[i] = x_i - x_i_minus_one;  
    }  
}
```

# Using shared memory

26

// Adjacent Difference application:

```
// Compute result[i] = input[i] - input[i-1]
```

**How do we use device memory bandwidth?**

```
__global__ void adj_diff_naive(int *result, int *input) {
    // compute this thread's global index
    unsigned int i = blockDim.x * blockIdx.x + threadIdx.x;

    if(i > 0) {
        // each thread loads two elements from device memory
        int x_i = input[i];
        int x_i_minus_one = input[i-1];

        result[i] = x_i - x_i_minus_one;
    }
}
```

# Using shared memory

27

// Adjacent Difference application:  
// Compute  $result[i] = input[i] - input[i-1]$

```
__global__ void adj_diff_naive(int *result, int *input) {
    // compute this thread's global index
    unsigned int i = blockDim.x * blockIdx.x + threadIdx.x;

    if(i > 0) {
        // each thread loads two elements from device memory
        int x_i = input[i];
        int x_i_minus_one = input[i-1];

        result[i] = x_i - x_i_minus_one;
    }
}
```

**The next thread also reads  $input[i]$**

# Using shared memory

28

```
__global__ void adj_diff(int *result, int *input) {
    unsigned int i = blockDim.x * blockIdx.x + threadIdx.x;

    __shared__ int s_data[BLOCK_SIZE]; // shared, 1 elt / thread
    // each thread reads 1 device memory elt, stores it in s_data
    s_data[threadIdx.x] = input[i];

    // avoid race condition: ensure all loads are complete
    __syncthreads();

    if(threadIdx.x > 0) {
        result[i] = s_data[threadIdx.x] - s_data[threadIdx.x-1];
    } else if(i > 0) {
        // I am thread 0 in this block: handle thread block boundary
        result[i] = s_data[threadIdx.x] - input[i-1];
    }
}
```

# Using shared memory: coalescing

29

```
__global__ void adj_diff(int *result, int *input) {
    unsigned int i = blockDim.x * blockIdx.x + threadIdx.x;

    __shared__ int s_data[BLOCK_SIZE]; // shared, 1 elt / thread
    // each thread reads 1 device memory elt, stores it in s_data
    s_data[threadIdx.x] = input[i]; // COALESCED ACCESS!

    // avoid race condition: ensure all loads are complete
    __syncthreads();

    if(threadIdx.x > 0) {
        result[i] = s_data[threadIdx.x] - s_data[threadIdx.x-1];
    } else if(i > 0) {
        // I am thread 0 in this block: handle thread block boundary
        result[i] = s_data[threadIdx.x] - input[i-1];
    }
}
```

# A Common Programming Strategy

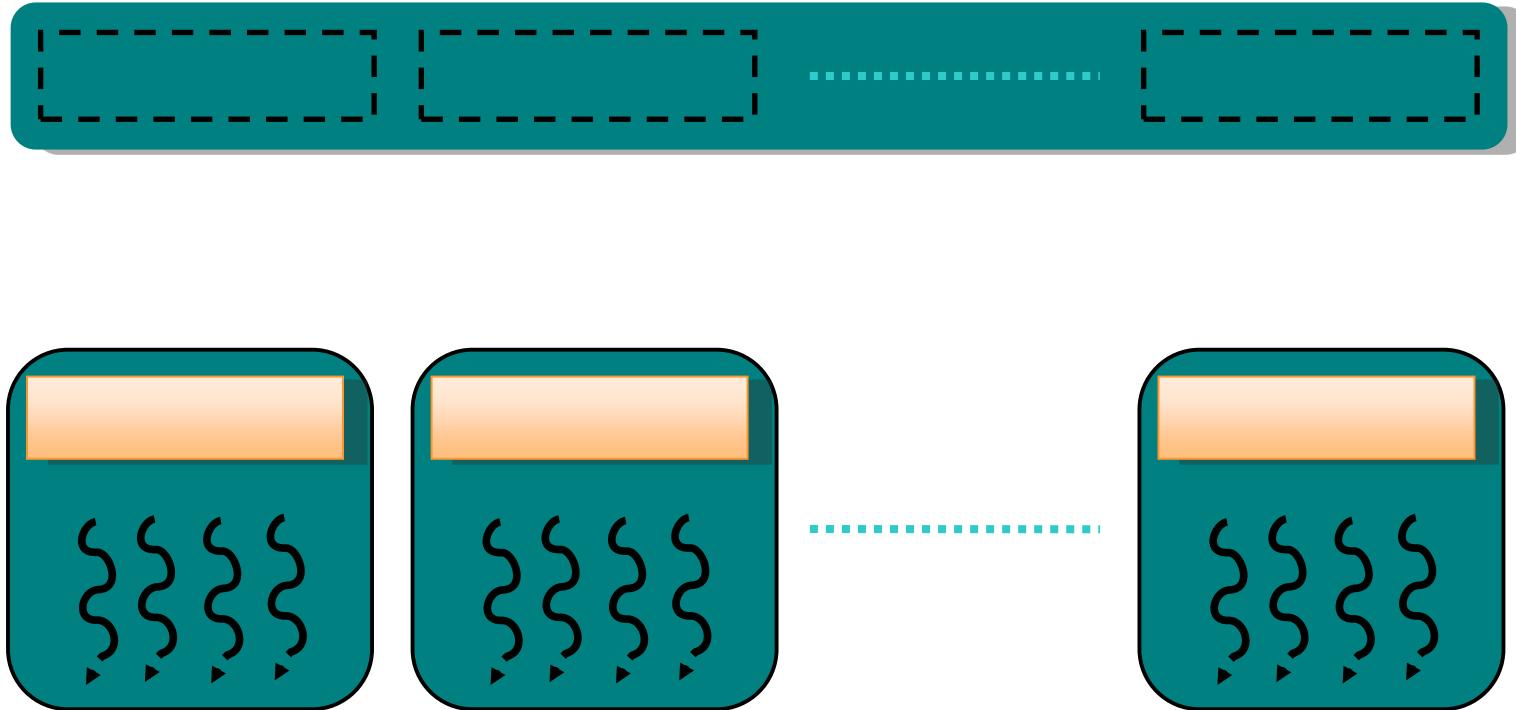
30



- Partition data into subsets that fit into shared memory

# A Common Programming Strategy

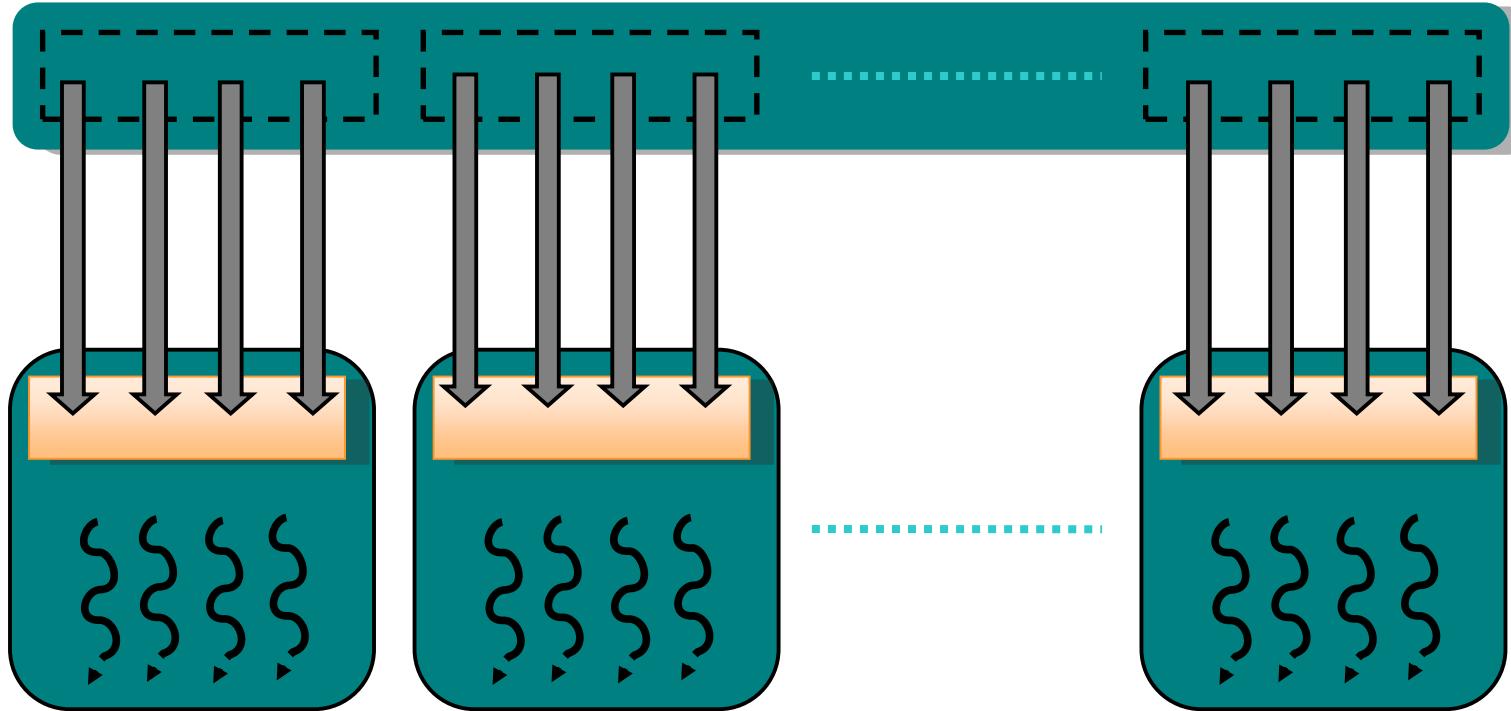
31



- Handle each data subset with one thread block

# A Common Programming Strategy

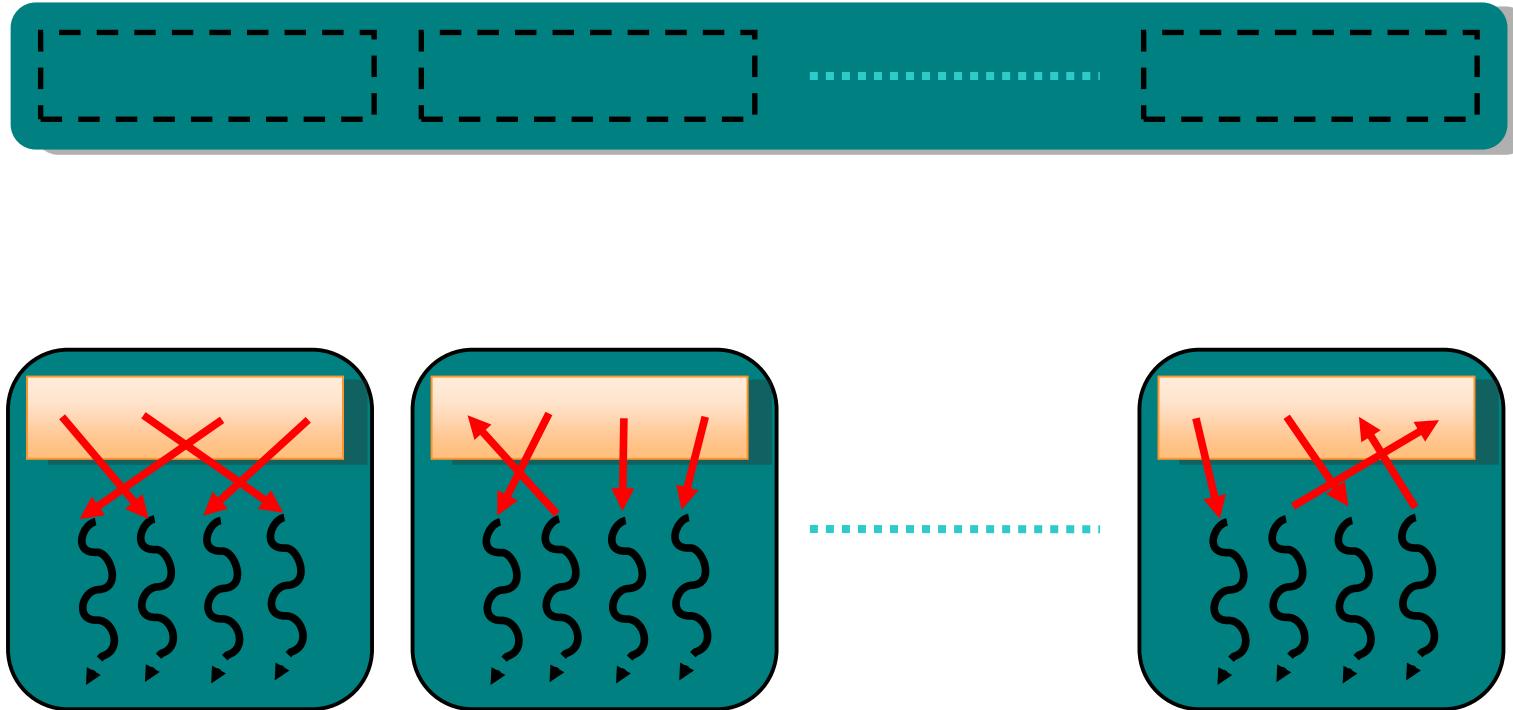
32



- Load the subset from device memory to shared memory, using multiple threads to exploit memory-level parallelism

# A Common Programming Strategy

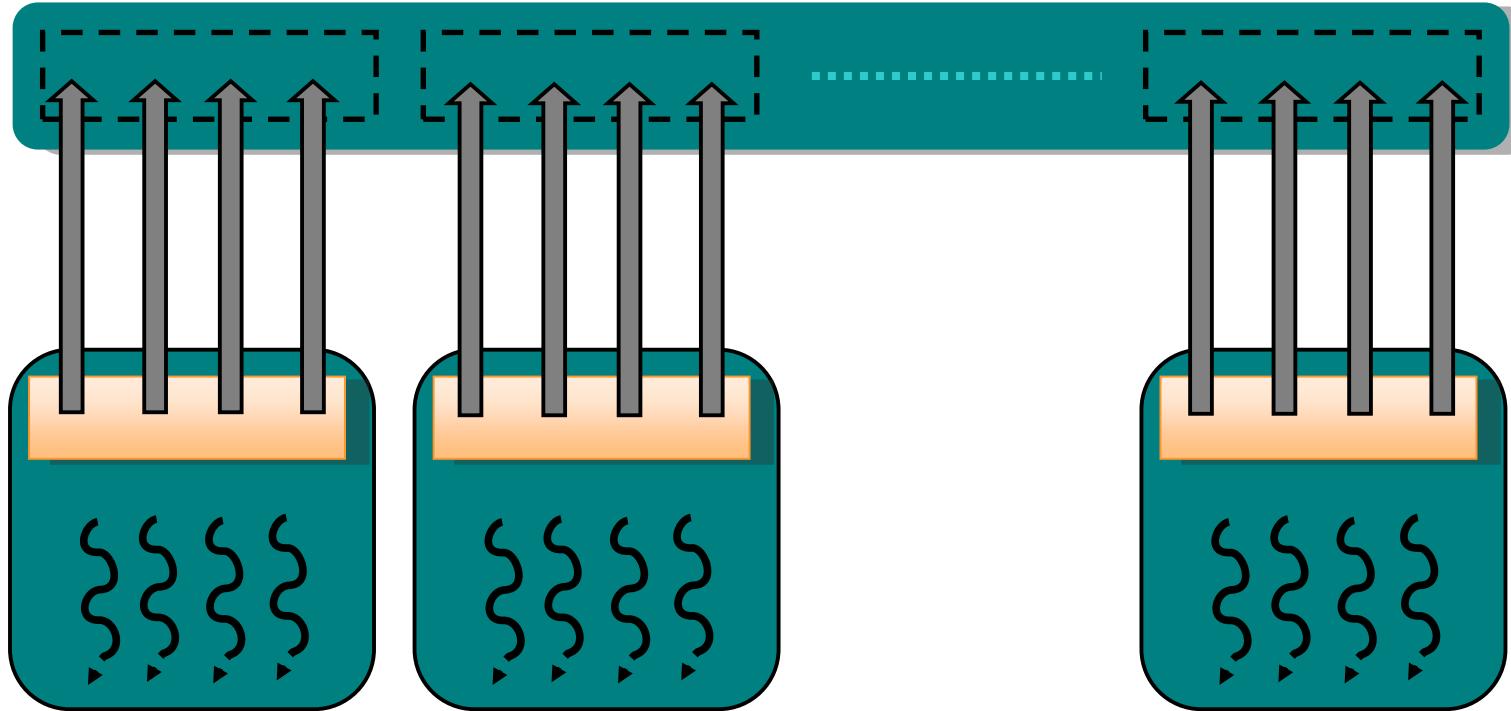
33



- Perform the computation on the subset from shared memory

# A Common Programming Strategy

34



- Copy the result from shared memory back to device memory

# CUDA: optimizing your application

## Optimizing Occupancy

# Thread Scheduling

36

- SM implements zero-overhead warp scheduling
  - A warp is a group of 32 threads that runs concurrently on a SM
  - At any time, only one of the warps is executed by SM
  - Warps whose next instruction has its inputs ready for consumption are eligible for execution
  - Eligible Warps are selected for execution on a prioritized scheduling policy
  - All threads in a warp execute the same instruction when selected

# Stalling warps

37

- What happens if all warps are stalled?
  - No instruction issued → performance lost
- Most common reason for stalling?
  - Waiting on global memory
- If your code reads global memory every couple of instructions
  - You should try to maximize occupancy

# Occupancy

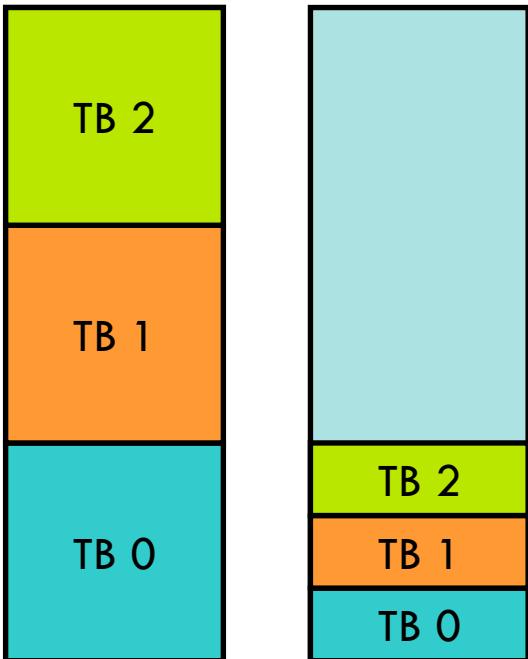
38

- What determines occupancy?
- Limited resources!
  - Register usage per thread
  - Shared memory per thread block

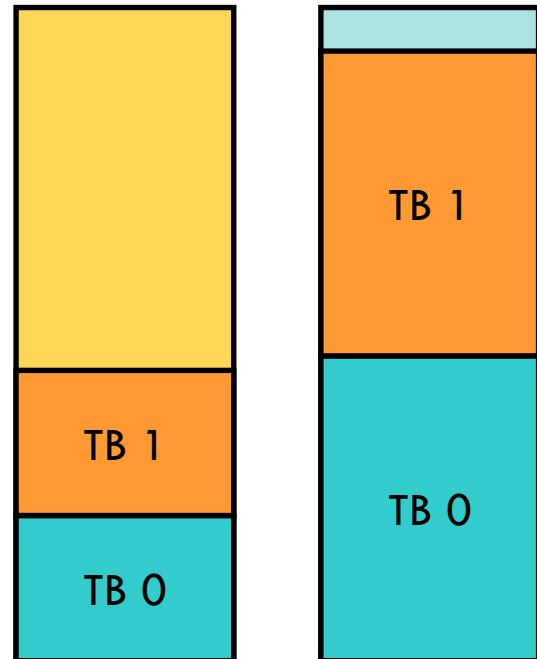
# Resource Limits (1)

39

Registers      Shared Memory



Registers      Shared Memory



- Pool of registers and shared memory per SM
  - Each thread block grabs registers & shared memory
  - If one or the other is fully utilized → no more thread blocks

# Resource Limits (2)

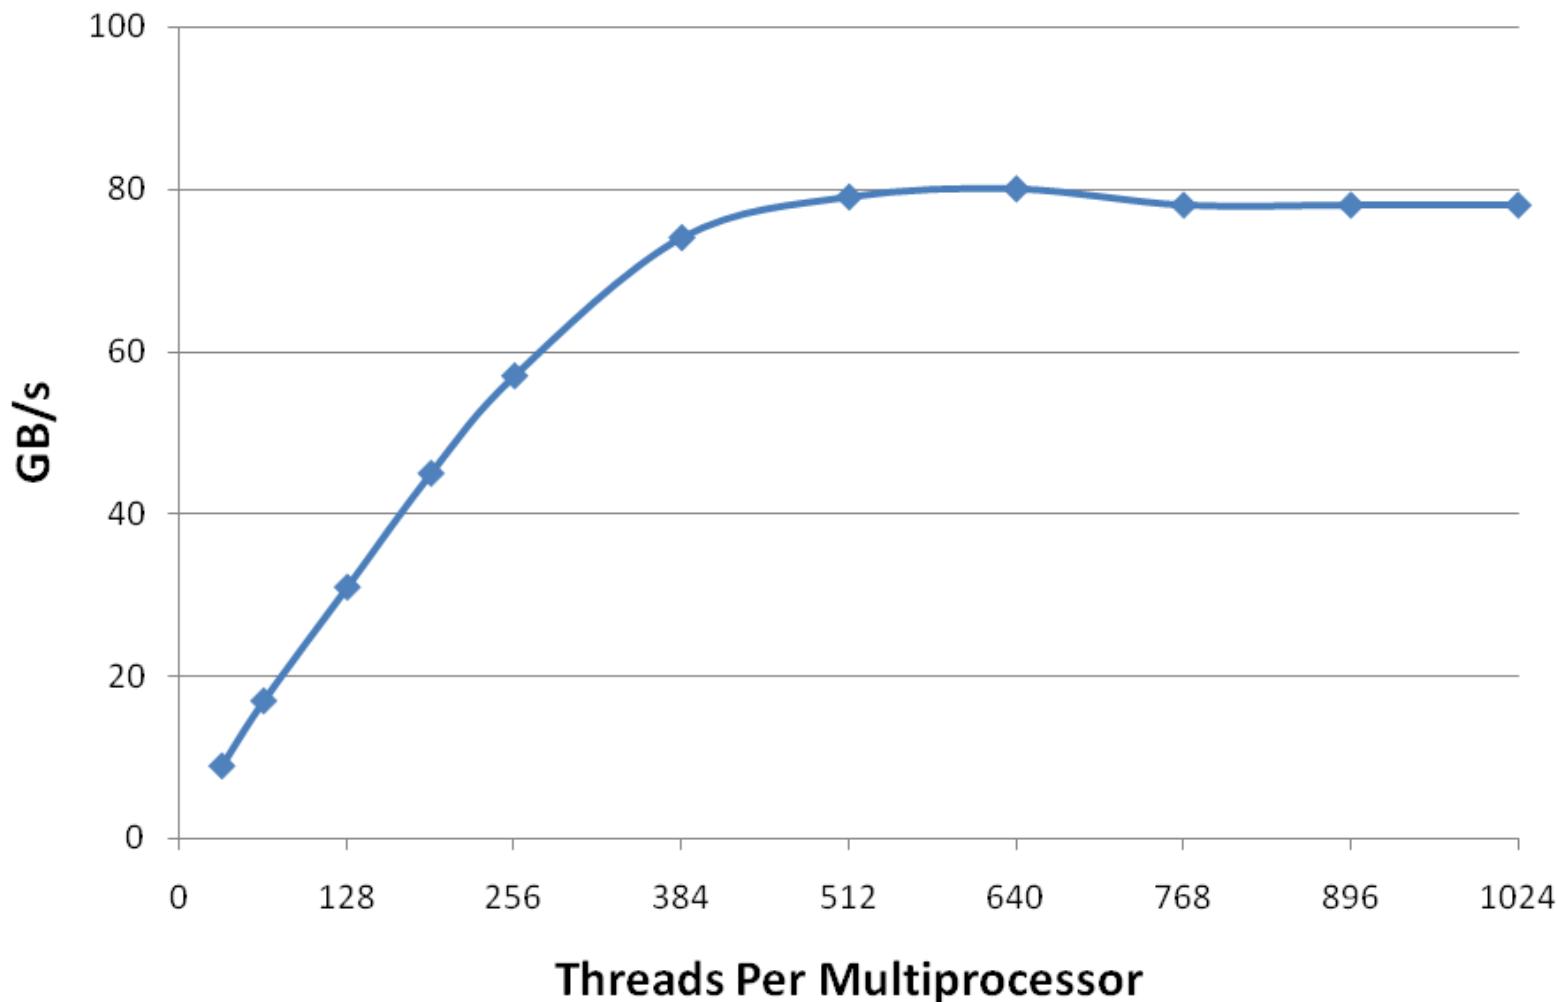
40

- Can only have a limited number of blocks per SM
  - If they're too small, can't fill up the SM
- Higher occupancy has diminishing returns for hiding latency

# Hiding Latency with more threads

41

Throughput, 32-bit words



# How do you know what you're using?

42

- Use “`--ptxas-options="-v"`” to get register and shared memory usage
- You can plug those numbers into CUDA Occupancy Calculator

`cmem[0]`:kernel arguments

`cmem[3]`:user defined constant objects

`cmem[16]`:compiler generated constants (some of which may correspond to literal constants in the source code)

# CUDA: optimizing your application

## Shared memory bank conflicts

# Shared Memory Banks

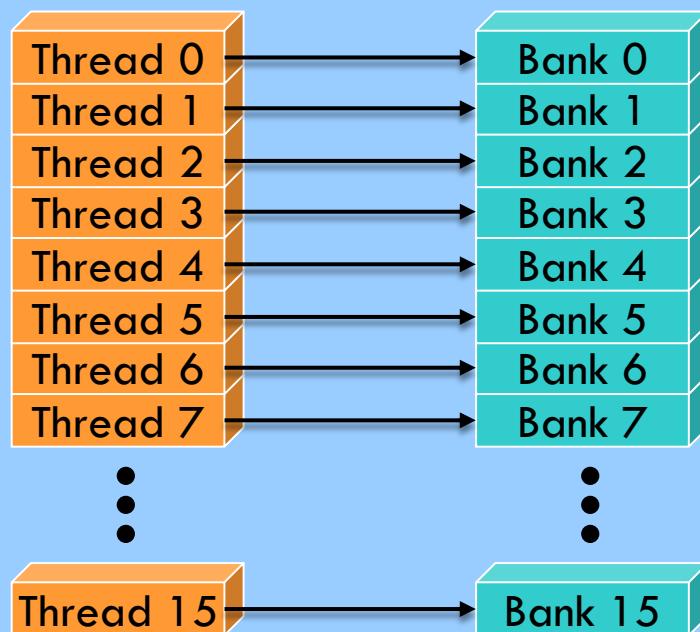
45

- Shared memory is banked
  - Only matters for threads within a warp
  - Full performance with some restrictions
    - Threads can each access different banks
    - Or can all access the same value
- Consecutive words are in different banks
- If two or more threads access the same bank but different value, we get bank conflicts

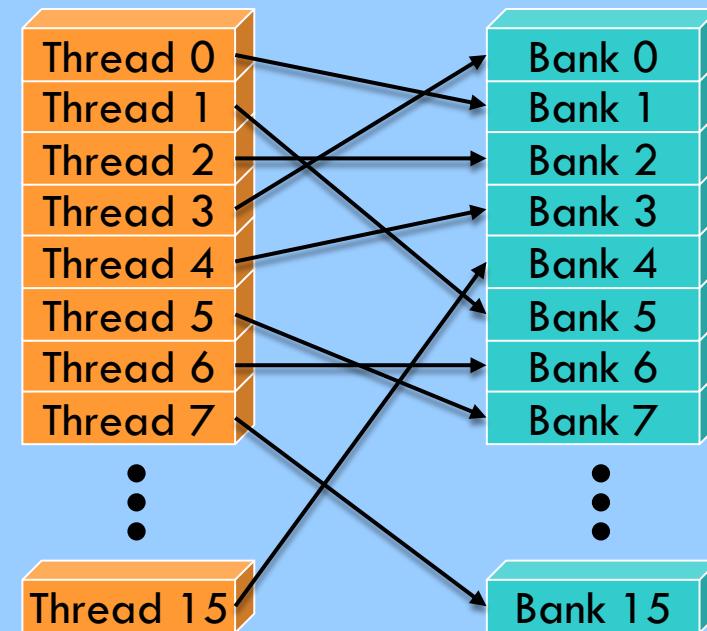
# Bank Addressing Examples: OK

46

- **No Bank Conflicts**



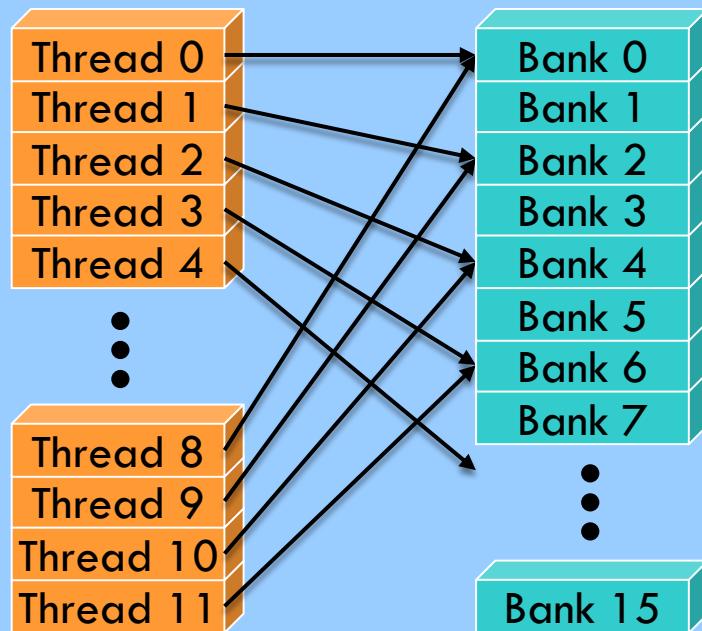
- **No Bank Conflicts**



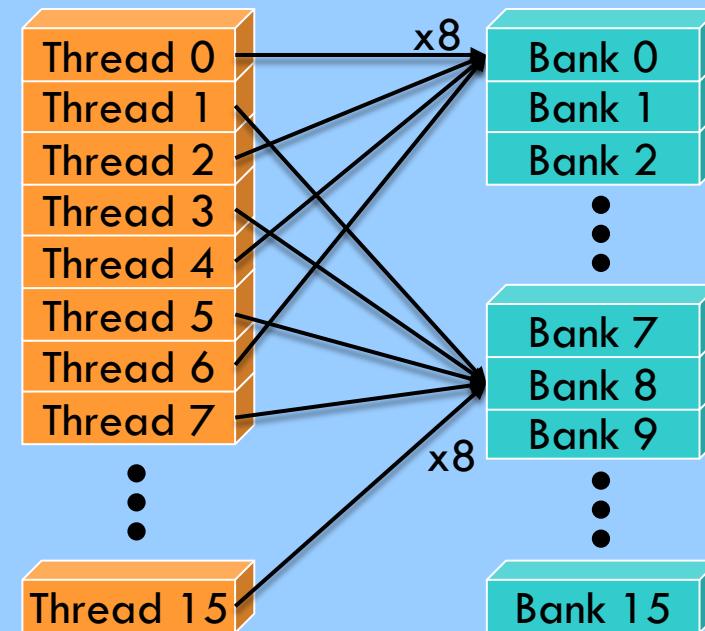
# Bank Addressing Examples: BAD

47

- 2-way Bank Conflicts



- 8-way Bank Conflicts



# Trick to Assess Performance Impact

48

- Change all shared memory reads to the same value
- All broadcasts = no conflicts
- Will show how much performance could be improved by eliminating bank conflicts
- The same doesn't work for shared memory writes
  - So, replace shared memory array indices with **threadIdx.x**
  - (Could also be done for the reads)

# Other useful tools

cuda-memcheck

nvprof

# Summary and conclusions

62

- Higher performance cannot be reached by increasing clock frequencies anymore
- Solution: introduction of large-scale parallelism
- Multiple cores on a chip
  - Today:
    - Up to 100 CPU cores in a node
    - Up to 5000 cores on a single GPU
    - Host system can contain multiple GPUs: 10,000+ cores
    - We can build clusters of these nodes!
  - Future: 100,000s – millions of cores?

# Summary and conclusions

63

- Many different types of many-core hardware
- Very different properties
  - Performance
  - Programmability
  - Portability
- It's all about the memory
- Choose the right platform for your application
  - Arithmetic intensity / Operational intensity
  - Roofline model