


# UNIFIED MEMORY

# UNIFIED MEMORY

Dramatically Lower Developer Effort

CUDA 6+



Simpler  
Programming &  
Memory Model

Single allocation, single pointer,  
accessible anywhere  
Eliminate need for *explicit copy*  
Greatly simplifies code porting

Performance  
Through  
Data Locality

Migrate data to accessing processor  
Guarantee global coherence  
Still allows explicit hand tuning

# SIMPLIFIED MEMORY MANAGEMENT CODE

## CPU Code

```
void sortfile(FILE *fp, int N) {  
    char *data;  
    data = (char *)malloc(N);  
  
    fread(data, 1, N, fp);  
  
    qsort(data, N, 1, compare);  
  
    use_data(data);  
  
    free(data);  
}
```

## CUDA 6 Code with Unified Memory

```
void sortfile(FILE *fp, int N) {  
    char *data;  
    cudaMallocManaged(&data, N);  
  
    fread(data, 1, N, fp);  
  
    qsort<<<...>>>(data,N,1,compare);  
    cudaDeviceSynchronize();  
  
    use_data(data);  
  
    cudaFree(data);  
}
```

# GREAT PERFORMANCE WITH UNIFIED MEMORY


RAJA: Portable C++ Framework for parallel-for style programming

RAJA uses Unified Memory for heterogeneous array allocations

Parallel forall loops run on device

“Excellent performance considering this is a “generic” version of LULESH with no architecture-specific tuning.”


-Jeff Keasler, LLNL



GPU: NVIDIA Tesla K40, CPU: Intel Haswell E5-2650 v3 @ 2.30GHz, single socket 10-core

# CUDA 8: UNIFIED MEMORY

Large datasets, simple programming, High Performance



# UNIFIED MEMORY EXAMPLE

## On-Demand Paging

```
__global__
void setvalue(int *ptr, int index, int val)
{
    ptr[index] = val;
}
```

```
void foo(int size) {
    char *data;
    cudaMallocManaged(&data, size);
    memset(data, 0, size);
    setvalue<<<...>>>(data, size/2, 5);
    cudaDeviceSynchronize();
    useData(data);
}
```



Unified Memory allocation



Access all values on CPU



Access one value on GPU


# HOW UNIFIED MEMORY WORKS IN CUDA 6

## Servicing CPU page faults

GPU Code

```
__global__  
void setvalue(char *ptr, int index, char val)  
{  
    ptr[index] = val;  
}
```


GPU Memory Mapping



CPU Code

```
cudaMallocManaged(&array, size);  
memset(array, size);  
setvalue<<<...>>>(array, size/2, 5);
```

CPU Memory Mapping



Page Fault

Interconnect


# HOW UNIFIED MEMORY WORKS ON PASCAL

## Servicing CPU *and* GPU Page Faults

GPU Code

```
__global__  
void setvalue(char *ptr, int index, char val)  
{  
    ptr[index] = val;  
}
```


GPU Memory Mapping



CPU Code


```
cudaMallocManaged(&array, size);  
memset(array, size);  
setvalue<<<...>>>(array, size/2, 5);
```

CPU Memory Mapping



# USE CASE: ON-DEMAND PAGING


## Graph Algorithms



Performance over GPU directly accessing host memory  
(zero-copy)

Unified Memory speed-up over system memory

baseline    optimized



Baseline: migrate on first touch  
Optimized: best placement in memory

# UNIFIED MEMORY ON PASCAL

## GPU memory oversubscription

```
void foo() {  
    // Assume GPU has 16 GB memory  
    // Allocate 32 GB  
    char *data;  
    size_t size = 32*1024*1024*1024;  
    cudaMallocManaged(&data, size);  
}
```

32 GB allocation

Pascal supports allocations where only a subset of pages reside on GPU. Pages can be migrated to the GPU when “hot”.

Fails on Kepler/Maxwell

# GPU OVERSUBSCRIPTION


## Now possible with Pascal

Many domains would benefit from GPU memory oversubscription:

**Combustion** - many species to solve for

**Quantum chemistry** - larger systems

**Ray tracing** - larger scenes to render



# UNIFIED MEMORY ON PASCAL

Concurrent CPU/GPU access to managed memory

```
__global__ void mykernel(char *data) {
    data[1] = 'g';
}

void foo() {
    char *data;
    cudaMallocManaged(&data, 2);

    mykernel<<<...>>>(data);
    // no synchronize here
    data[0] = 'c';

    cudaFree(data);
}
```


OK on Pascal: just a page fault

Concurrent CPU access to ‘data’ on previous  
GPUs caused a fatal segmentation fault

# UNIFIED MEMORY ON PASCAL

## System-Wide Atomics

```
__global__ void mykernel(int *addr) {  
    atomicAdd(addr, 10);  
}  
  
void foo() {  
    int *addr;  
    cudaMallocManaged(addr, 4);  
    *addr = 0;  
  
    mykernel<<<...>>>(addr);  
    __sync_fetch_and_add(addr, 10);  
}
```



- Pascal enables system-wide atomics
- Direct support of atomics over NVLink
  - Software-assisted over PCIe

System-wide atomics not available on Kepler / Maxwell

# PERFORMANCE TUNING ON PASCAL

## Explicit Memory Hints and Prefetching

Advise runtime on known memory access behaviors with `cudaMemAdvise()`

`cudaMemAdviseSetReadMostly`: Specify read duplication

`cudaMemAdviseSetPreferredLocation`: suggest best location

`cudaMemAdviseSetAccessedBy`: initialize a mapping

Explicit prefetching with `cudaMemPrefetchAsync(ptr, length, destDevice, stream)`

Unified Memory alternative to `cudaMemcpyAsync`

Asynchronous operation that follows CUDA stream semantics

To Learn More:  
S6216 “The Future of Unified Memory” by Nikolay Sakharnykh  
Tuesday, 4pm