


## CUDA STREAMS

BEST PRACTICES AND COMMON PITFALLS

Justin Luitjens - NVIDIA



# Simple Processing Flow




# CONCURRENCY THROUGH PIPELINING

- Serial




- Concurrent- overlap kernel and D2H copy




# CONCURRENCY THROUGH PIPELINING


- Serial (1x)




- 2-way concurrency (up to 2x)




- 3-way concurrency (up to 3x)



- **4-way concurrency (3x+)**



- **4+ way concurrency**




# EXAMPLE - TILED DGEMM

- CPU (dual 6 core SandyBridge E5-2667 @2.9 Ghz, MKL)
  - 222 Gflop/s
- GPU (K20X)
  - Serial: 519 Gflop/s (2.3x)
  - 2-way: 663 Gflop/s (3x)
  - 3-way: 990 Gflop/s (4x)
- GPU + CPU
  - 4-way con.: 1180 Gflop/s (5.3x)
- Obtain maximum performance by leveraging concurrency
- All PCI-E traffic is hidden
  - Effectively removes device memory size limitations!

DGEMM: m=n=16384, k=1408


Nvidia Visual Profiler (nvvp)



# ENABLING CONCURRENCY WITH STREAMS

# SYNCHRONICITY IN CUDA

- All CUDA calls are either synchronous or asynchronous w.r.t the host
  - **Synchronous**: enqueue work and wait for completion
  - **Asynchronous**: enqueue work and return immediately
- Kernel Launches are asynchronous Automatic overlap with host



# CUDA STREAMS

- A **stream** is a queue of device work
  - The host places work in the queue and continues on immediately
  - Device schedules work from streams when resources are free
- CUDA operations are placed within a stream
  - e.g. Kernel launches, memory copies
- Operations within the **same stream** are **ordered** (FIFO) and cannot overlap
- Operations in **different streams** are **unordered** and can overlap

# MANAGING STREAMS

- **`cudaStream_t stream;`**
  - Declares a stream handle
- **`cudaStreamCreate(&stream);`**
  - Allocates a stream
- **`cudaStreamDestroy(stream);`**
  - Deallocates a stream
  - Synchronizes host until work in stream has completed

# PLACING WORK INTO A STREAM

- Stream is the 4<sup>th</sup> launch parameter
  - `kernel<<< blocks , threads, smem, stream>>>();`
- Stream is passed into some API calls
  - `cudaMemcpyAsync( dst, src, size, dir, stream);`

# DEFAULT STREAM


- Unless otherwise specified all calls are placed into a default stream
  - Often referred to as “Stream 0”
- Stream 0 has special synchronization rules
  - Synchronous with all streams
    - Operations in stream 0 cannot overlap other streams
- Exception: Streams with non-blocking flag set

```
cudaStreamCreateWithFlags(&stream, cudaStreamNonBlocking)
```

# KERNEL CONCURRENCY


- Assume foo only utilizes 50% of the GPU
- Default stream

```
foo<<<blocks, threads>>>();  
foo<<<blocks, threads>>>();
```



- Default & user streams


```
cudaStream_t stream1;  
cudaStreamCreate(&stream1);  
foo<<<blocks, threads>>>();  
foo<<<blocks, threads, 0, stream1>>>();  
cudaStreamDestroy(stream1);
```



# KERNEL CONCURRENCY

- Assume foo only utilizes 50% of the GPU
- Default & user streams


```
cudaStream_t stream1;  
  
cudaStreamCreateWithFlags(&stream1, cudaStreamNonBlocking);  
  
foo<<<blocks, threads>>>();  
  
foo<<<blocks, threads, 0, stream1>>>();  
  
cudaStreamDestroy(stream1);
```



# KERNEL CONCURRENCY

- Assume foo only utilizes 50% of the GPU  
User streams

```
cudaStream_t stream1, stream2;  
  
cudaStreamCreate(&stream1);  
cudaStreamCreate(&stream2);  
foo<<<blocks, threads, 0, stream1>>>();  
foo<<<blocks, threads, 0, stream2>>>();  
cudaStreamDestroy(stream1);  
cudaStreamDestroy(stream2);
```



# REVIEW

- The host is automatically asynchronous with kernel launches
- Use streams to control asynchronous behavior
  - Ordered within a stream (FIFO)
  - Unordered with other streams
  - Default stream is synchronous with all streams.

# Concurrent Memory Copies

# CONCURRENT MEMORY COPIES

- First we must review CUDA memory

# THREE TYPES OF MEMORY

- Device Memory
  - Allocated using `cudaMalloc`
  - Cannot be paged
- Pageable Host Memory
  - Default allocation (e.g. `malloc`, `calloc`, `new`, etc)
  - Can be paged in and out by the OS
- Pinned (Page-Locked) Host Memory
  - Allocated using special allocators
  - Cannot be paged out by the OS

# ALLOCATING PINNED MEMORY

- **cudaMallocHost(...)** / **cudaHostAlloc(...)**
  - Allocate/Free pinned memory on the host
  - Replaces malloc/free/new
- **cudaFreeHost(...)**
  - Frees memory allocated by cudaMallocHost or cudaHostAlloc
- **cudaHostRegister(...)** / **cudaHostUnregister(...)**
  - Pins/Unpins pageable memory (making it pinned memory)
  - Slow so don't do often
- Why pin memory?
  - Pagable memory is transferred using the host CPU
  - Pinned memory is transferred using the DMA engines
    - Frees the CPU for asynchronous execution
    - Achieves a higher percent of peak bandwidth

# CONCURRENT MEMORY COPIES

- **cudaMemcpy( . . . )**
  - Places transfer into default stream
  - Synchronous: Must complete prior to returning
- **cudaMemcpyAsync( . . . , &stream)**
  - Places transfer into stream and returns immediately
- To achieve concurrency
  - Transfers must be in a non-default stream
  - Must use async copies
  - 1 transfer per direction at a time
  - Memory on the host must be pinned

# PAGED MEMORY EXAMPLE

```
int *h_ptr, *d_ptr;  
  
h_ptr=malloc(bytes);  
cudaMalloc(&d_ptr,bytes);  
  
cudaMemcpy(d_ptr,h_ptr,bytes,cudaMemcpyHostToDevice);  
  
free(h_ptr);  
cudaFree(d_ptr);
```

# PINNED MEMORY: EXAMPLE 1

```
int *h_ptr, *d_ptr;  
  
cudaMallocHost(&h_ptr,bytes);  
cudaMalloc(&d_ptr,bytes);  
  
cudaMemcpy(d_ptr,h_ptr,bytes,cudaMemcpyHostToDevice);  
  
cudaFreeHost(h_ptr);  
cudaFree(d_ptr);
```

# PINNED MEMORY: EXAMPLE 2

```
int *h_ptr, *d_ptr;  
  
h_ptr=malloc(bytes);  
cudaHostRegister(h_ptr,bytes,0);  
cudaMalloc(&d_ptr,bytes);  
  
cudaMemcpy(d_ptr,h_ptr,bytes,cudaMemcpyHostToDevice);  
  
cudaHostUnregister(h_ptr);  
free(h_ptr);  
cudaFree(d_ptr);
```

# CONCURRENCY EXAMPLES

## Synchronous



```
cudaMemcpy(...);  
foo<<<...>>>();
```

## Asynchronous Same Stream

```
cudaMemcpyAsync(..., stream1);  
foo<<<..., stream1>>>();
```

## Asynchronous Different Streams

```
cudaMemcpyAsync(..., stream1);  
foo<<<..., stream2>>>();
```



# REVIEW

- Memory copies can execute concurrently if (and only if)
  - The memory copy is in a different non-default stream
  - The copy uses pinned memory on the host
  - The asynchronous API is called
  - There isn't another memory copy occurring in the same direction at the same time.

# Synchronization

# SYNCHRONIZATION APIs

- Synchronize everything
  - `cudaDeviceSynchronize()`
    - Blocks host until all issued CUDA calls are complete
- Synchronize host w.r.t. a specific stream
  - `cudaStreamSynchronize( stream)`
    - Blocks host until all issued CUDA calls in stream are complete
- Synchronize host or devices using events

More  
Synchronization



Less  
Synchronization

# CUDA EVENTS

- Provide a mechanism to signal when operations have occurred in a stream
  - Useful for profiling and synchronization
- Events have a boolean state:
  - Occurred
  - Not Occurred
  - **Important:** Default state = occurred

# MANAGING EVENTS

- **cudaEventCreate (&event)**
  - Creates an event
- **cudaEventDestroy (&event)**
  - Destroys an event
- **cudaEventCreateWithFlags (&ev, cudaEventDisableTiming)**
  - Disables timing to increase performance and avoid synchronization issues
- **cudaEventRecord (&event, stream)**
  - Set the event state to not occurred
  - Enqueue the event into a stream
  - Event state is set to occurred when it reaches the front of the stream

# SYNCHRONIZATION USING EVENTS

- Synchronize using events
  - `cudaEventQuery ( event )`
    - Returns CUDA\_SUCCESS if an event has occurred
  - `cudaEventSynchronize ( event )`
    - Blocks host until stream completes all outstanding calls
  - `cudaStreamWaitEvent ( stream, event )`
    - Blocks stream until event occurs
    - Only blocks launches after this call
    - Does not block the host!
- Common multi-threading mistake:
  - Calling `cudaEventSynchronize` before `cudaEventRecord`

# CUDA\_LAUNCH\_BLOCKING

- Environment variable which forces synchronization
  - `export CUDA_LAUNCH_BLOCKING=1`
  - All CUDA operations are synchronous w.r.t the host
- Useful for debugging race conditions
  - If it runs successfully with `CUDA_LAUNCH_BLOCKING` set but doesn't without you have a race condition.

# REVIEW

- Synchronization with the host can be accomplished via
  - `cudaDeviceSynchronize()`
  - `cudaStreamSynchronize(stream)`
  - `cudaEventSynchronize(event)`
- Synchronization between streams can be accomplished with
  - `cudaStreamWaitEvent(stream,event)`
- Use `CUDA_LAUNCH_BLOCKING` to identify race conditions


# COMMON STREAMING PROBLEMS

# COMMON STREAMING PROBLEMS

- The following is an attempt to demonstrate the most common streaming issues I've seen in customers applications
- They are loosely ordered according to how common they are

# CASE STUDY 1-A

```
for(int i=0;i<repeat;i++)  
{  
    kernel<<<1,1,0,stream1>>>();  
    kernel<<<1,1>>>();  
}
```




**Problem:**

**One kernel is in the default stream**

**Stream 2 is the default stream**

## CASE STUDY 1-A

```
for(int i=0;i<repeat;i++) {  
    kernel<<<1,1,0,stream1>>>();  
    kernel<<<1,1,0,stream2>>>();  
}
```



**Solution:**  
**Place each kernel in its own stream**


# PROBLEM 1: USING THE DEFAULT STREAM

- Symptoms
  - One stream will not overlap other streams
    - In Cuda 5.0 stream 2 = default stream
  - Search for `cudaEventRecord(event)` , `cudaMemcpyAsync()`, etc.
    - If stream is not specified it is placed into the default stream
  - Search for kernel launches in the default stream
    - `<<<a,b>>>`
- Solutions
  - Move work into a non-default stream
  - `cudaEventRecord(event,stream)`, `cudaMemcpyAsync(...,stream)`
  - Alternative: Allocate other streams as non-blocking streams

# CASE STUDY 2-A


```
for(int i=0;i<repeat;i++) {  
    cudaMemcpy(d_ptr,h_ptr,bytes, cudaMemcpyHostToDevice) ;  
    kernel<<<1,1,0,stream2>>>();  
    cudaDeviceSynchronize() ;  
}
```

**Problem:**  
**Memory copy is synchronous**



# CASE STUDY 2-A

```
for(int i=0;i<repeat;i++) {  
    cudaMemcpyAsync(d_ptr,h_ptr,bytes, cudaMemcpyHostToDevice, stream1);  
    kernel<<<1,1,0,stream2>>>();  
    cudaDeviceSynchronize();  
}
```




**Solution:**  
**Use asynchronous API**

# CASE STUDY 2-B


```
for(int i=0;i<repeat;i++) {  
    cudaMemcpyAsync(d_ptr,h_ptr,bytes, cudaMemcpyHostToDevice, stream1);  
    kernel<<<1,1,0,stream2>>>();  
    cudaDeviceSynchronize();  
}
```

Problem: ??



# CASE STUDY 2-B


```
for(int i=0;i<repeat;i++) {  
    cudaMemcpyAsync(d_ptr,h_ptr,bytes, cudaMemcpyHostToDevice, stream1);  
    kernel<<<1,1,0,stream2>>>();  
    cudaDeviceSynchronize();  
}
```



Host doesn't get ahead  
Cuda 5.5 reports "Pageable" type

# CASE STUDY 2-B

```
cudaHostRegister(h_ptr,bytes,0);
for(int i=0;i<repeat;i++) {
    cudaMemcpyAsync(d_ptr,h_ptr,bytes, cudaMemcpyHostToDevice, stream1);
    kernel<<<1,1,0,stream2>>>();
    cudaDeviceSynchronize();
}
cudaHostUnregister(h_ptr);
```



Solution:

Pin host memory using `cudaHostRegister` or `cudaMallocHost`

# PROBLEM 2: MEMORY TRANSFERS ISSUES

- Symptoms


- Memory copies do not overlap
- Host spends excessive time in memory copy API
- Cuda reports “Pageable” memory (Cuda 5.5+)

- Solutions

- Use asynchronous memory copies
- Use pinned memory for host memory
  - `cudaMallocHost` or `cudaHostRegister`

# CASE STUDY 3

```
void launchwork(cudaStream_t stream) {  
    int *mem;  
    cudaMalloc(&mem,bytes);  
    kernel<<<1,1,0,stream>>>(mem);  
    cudaFree(mem);  
}  
...  
  
for(int i=0;i<repeat;i++) {  
    launchwork(stream1);  
    launchwork(stream2);  
}
```




Problem:

Allocation & deallocation synchronize the device

Host blocked in allocation/free

# CASE STUDY 3

```
void launchwork(cudaStream_t stream, int *mem) {  
    kernel<<<1,1,0,stream>>>(mem);  
}  
...  
  
for(int i=0;i<repeat;i++) {  
    launchwork<1>(stream1,mem1);  
    launchwork<2>(stream2,mem2);  
}
```



**Solution:**


Separate CUDA memory and objects in streams and events

# PROBLEM 3: IMPLICIT SYNCHRONIZATION

- Symptoms
  - Host does not get ahead
  - Host shows excessive time in certain API calls
    - `cudaMalloc`, `cudaFree`, `cudaEventCreate`, `cudaEventDestroy`, `cudaStreamCreate`,  
`cudaStreamCreate`, `cudaHostRegister`, `cudaHostUnregister`,  
`cudaFuncSetCacheConfig`
- Solution:
  - Reuse memory and data structures

# CASE STUDY 4

```
for(int i=0;i<repeat;i++)  
{  
    hostwork();  
    kernel<<<1,1,0,stream1>>>();  
    hostwork();  
    kernel<<<1,1,0,stream2>>>();  
}
```



**Problem:**  
**Host is limiting performance**

Host is outside of API calls

## PROBLEM 4: LIMITED BY HOST

- Symptoms
  - Host is outside of cuda APIs
  - Large gaps in timeline where the host and device are empty
- Solution
  - Move more work to the GPU
  - Multi-thread host code

# PROBLEM 5: LIMITED BY LAUNCH OVERHEAD

- Symptoms

- Host does not get ahead
  - Kernels are short <30 us
  - Time between successive kernels is >10 us

- Solutions

- Make longer running kernels
    - Fuse nearby kernels together
    - Batch work within a single kernel
    - Solve larger problems